

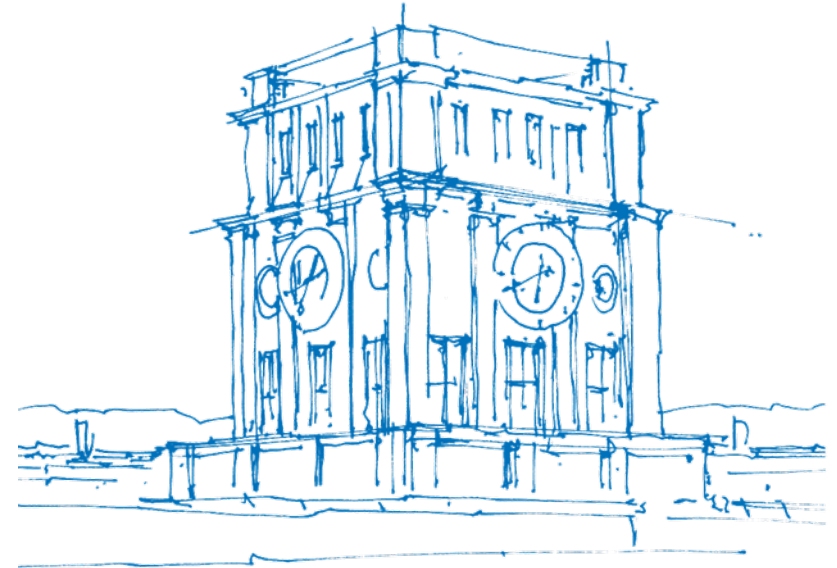
Parallel Programming Tutorial - Introduction to Pthread API

Bengisu Elis, M.Sc.

Chair for Computer Architecture and Parallel Systems (CAPS)

Technical University Munich

7. Mai 2019



TUM Uhrenturm

Organization

Organization

- Course web-page
 - parprog.caps.in.tum.de
 - Register and login using your LRZId's (@mytum IDs).
 - Course schedule; lecture and tutorial
 - Exercises and assignment submission
 - Lecture and Tutorial slides are on Moodle !
- Tutorials: Wednesdays at 8:15 ? to 9:45 ?
 - Always check the schedule in the web-page
- Where to find us?
 - Chair for Computer Architecture and Parallel Systems (Prof. Dr. M. Schulz)
 - My email address is: bengisu.elis@tum.de
 - Room: MI, 01.04.053

Parallel Programming

[Login](#)

[About](#)

[Schedule](#)

[Contacts](#)

[Impressum](#)

[Datenschutz](#)

Schedule (tentative)

Week #	Date	Description; lecture and tutorial topics (tentative)
1	29.04.2019	Lecture 01: Introduction / Basic approaches / Types of Parallelism / Metrics / Classification of Languages and Libraries
1	01.05.2019	Holiday : Maifeiertag
2	06.05.2019	Lecture 02: Threading (shared memory, general idea, Pthreads, User-level threads)
2	08.05.2019	Tutorial 01: Introduction to Pthreads API
3	13.05.2019	Tutorial 02: Pthread - Locks and Synchronization
3	15.05.2019	Lecture 03: OpenMP Basics
4	20.05.2019	Lecture 04: Dependency Analysis
4	22.05.2019	Tutorial 03: Parallelism in C++
5	27.05.2019	Lecture 05: OpenMP Advanced
5	29.05.2019	Tutorial 04: OpenMP Basics
6	03.06.2019	Lecture 06: HPC Architectures / Applications / SuperMUC / Using HPC resources
6	05.06.2019	Tutorial 05: OpenMPI Advanced
7	10.06.2019	Holiday : Pfingstmontag
7	12.06.2019	Lecture 07: MPI Basics
8	17.06.2019	Tutorial 06: Q&A on OpenMP

Assignments

You have 1 week to complete each assignment

- We will work on 10 assignments on parallel programming techniques
- Submission of 80% of the assignments brings you 0.3 bonus
- Submission server: <http://parprog.caps.in.tum.de>
 - Walk through of the submission work-flow at the end of today's tutorial session
- Submissions will be checked for:
 - Plagiarism, correctness (output, threads, synchronization), speedup, memory leaks
- Example solutions will be presented at the following tutorial session
- Topics
 - Pthreads (Posix Threads)
 - C++(11/14/17)
 - OpenMP (Open Multi-Processing)
 - Dependency analysis
 - MPI (Message Passing Interface)

Assistance on Assignments

Starting this week

Given by:

- Hasan Ashraf
hasan.ashraf@tum.de
- Philipp Czerner
philipp.czerner@tum.de

If you have questions regarding assignments and solutions, write an email to our tutors.

Resources

- POSIX Threads Programming
- An Introduction to Parallel Programming, by Peter Pacheco
- Programming with Posix Threads, by David Butenhof
- Patterns for Parallel Programming, by Timothy G. Mattson; Beverly A. Sanders; Berna L. Massingill
- Multithreading in Modern C++, by Rainer Grimm

Course Prerequisites

- knowledge of C/C++ (our code examples and assignments are all in C/C++)
 - memory management
 - pointers /references
 - global vs. static variables
- C/C++(11/14/17) books
 - (C89) The C Programming Language, Second Edition, by Brian W. Kernighan; Dennis M. Ritchie
 - (C99) C Primer Plus, Fifth Edition, by Stephen Prata
 - (C++11/14) The C++ Programming Language, Fourth Edition, by Bjarne Stroustrup
- Experience with Linux Command Line
- Resources
 - Book: The Linux Command Line
 - Basic video introduction: The Shell
- Knowing GCC
 - An Introduction to GCC, by Brian Gough

Posix Thread Programming

Posix Thread Programming

Definition: (Software) Thread

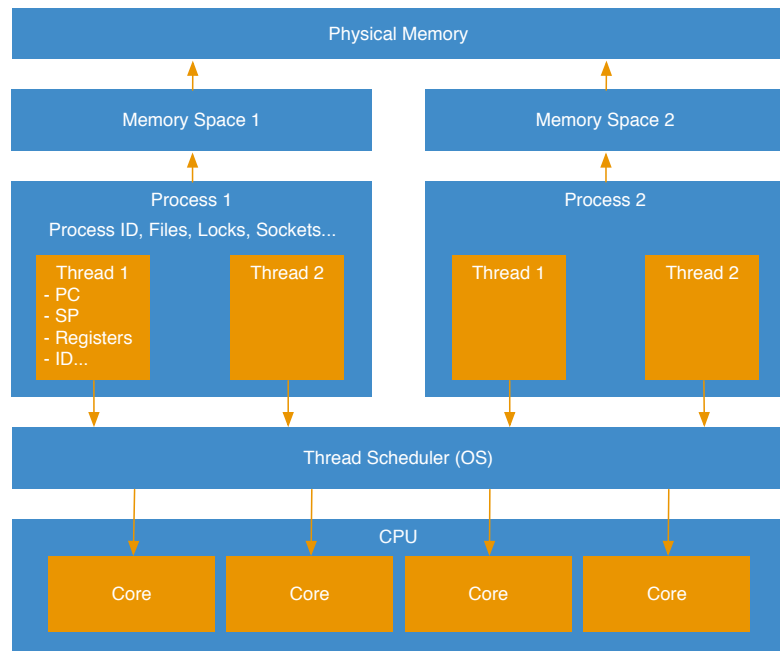
A thread is an independent stream of instructions that can be scheduled to run as such by the operating system. (Own PC and SP)

POSIX Threads (Pthreads)

- Were defined in 1995 (IEEE Std 1003.1c-1995)
- Is an API that defines a set of types, functions and constants
- Is implemented with a `pthread.h` header and a thread library
- Natively supported by FreeBSD, NetBSD, OpenBSD, Linux, Mac OS X, Android and Solaris
- Functions can be categorized in four groups:
 - Thread management
 - Mutexes
 - Condition variables

Why use Multithreading?

- **Performance gains**
Parallel processing by multiple processor cores
- **Increased application throughput**
Asynchronous system calls possible
- **Increased application responsiveness**
Application does not need to block operations
- **Replacing process-to-process communications**
Threads may communicate by shared-memory
- **Efficient use of system resources**
Lightweight context switches possible
- **Separation of concerns**
Some problems are inherently concurrent



Pthread Syntax / Semantics

Create Pthreads

```
1 int pthread_create(pthread_t *thread,  
2     const pthread_attr_t *attr,  
3     void *(*start_routine) (void *),  
4     void *arg);
```

- pthread_t *thread,
 - Pointer to thread identifier.
- const pthread_attr_t *attr
 - Optional pointer to pthread_attr_t to define behavior, if NULL defaults are used.
- void *(*start_routine) (void *)
 - Pointer to function prototype that is started. Function takes void pointer as argument and returns a void pointer.
- void *arg
 - Pointer to the argument that is used for the executed function.

Waiting for Pthread to finish

```
1 int pthread_join(pthread_t thread,  
2 void **retval);
```

- pthread_t thread,
 - Thread identifier, for which this function is waiting.
- void **retval
 - Optional pointer pointing to a void pointer. This can be used to return data of undefined size.

Example 1; creating a thread

```
1 #include <stdio.h>
2 #include <pthread.h>
3
4 // function to be executed by the thread
5 void* kernel (void* args){
6     printf("hello from the thread!\n");
7     return NULL;
8 }
9
10 int main(int argc, char *argv[])
11 {
12     pthread_t thread;                // allocate a thread
13     pthread_create(&thread, NULL, kernel , NULL); // create the thread and start executing kernel in parallel to main thread
14     printf("hello from main\n");
15     pthread_join(thread, NULL);      //wait for the thread to finish executing kernel
16
17     return 0;
18 }
```

Compile & Output

```
gcc -pthread -Wall -o ex1 ex1.c  
./ex1
```

```
Hello from main!  
Hello from the thread!
```

Example 2; creating multiple threads

```
1  int main(int argc, char *argv[])
2  {
3      //allocate the threads
4      int num_threads=4;
5      pthread_t *threads = (pthread_t*) malloc (num_threads *sizeof(pthread_t));
6
7      //create threads, start executing kernel in parallel
8      for (int i = 0; i < num_threads; ++i) {
9          pthread_create(&threads[i], NULL, kernel, NULL);
10     }
11
12     //wait for all the threads to finish executing kernel
13     for (int i = 0; i < num_threads; ++i) {
14         pthread_join(threads[i], NULL);
15     }
16
17     free(threads);
18     return 0;
19 }
```


Output

```
./ex2
```

```
Hello from the thread!  
Hello from the thread!  
Hello from the thread!  
Hello from the thread!
```

Example 3, passing an argument to threads

```
1 void* kernel (void* args){  
2     int id = *(int*)args;  
3     printf("Hello from the thread, myid: %d!\n", id);  
4     return NULL;  
5 }
```

Example 3, passing an argument to threads (cont.)

```
1  int main(int argc, char *argv[])
2  {
3      //allocate the threads
4      int num_threads=4;
5      pthread_t *threads = (pthread_t*) malloc (num_threads*sizeof(pthread_t));
6      int* id = (int*) malloc (num_threads*sizeof(int));
7
8      //create threads, start executing kernel in parallel
9      for (int i = 0; i < num_threads; ++i) {
10         id[i]=i; //set the id for the threads
11         pthread_create(&threads[i], NULL, kernel, id+i); //pass the id as argument to the threads
12     }
13
14     //wait for all the threads to finish executing kernel
15     for (int i = 0; i < num_threads; ++i) {
16         pthread_join(threads[i], NULL);
17     }
18
19     free(threads); free(id);
20     return 0;
21 }
```

Output

```
./ex3
```

```
Hello from the thread, myid: 1!  
Hello from the thread, myid: 0!  
Hello from the thread, myid: 2!  
Hello from the thread, myid: 3!
```

Example 4; process and thread IDs

```
1 void* kernel (void* args){
2     int id = *(int*)args;
3     printf("Hello from the thread, myid: %d, PID: %d, TID:%d!\n", id, getpid(), (int) gettid());
4     return NULL;
5 }
```

Output

```
./ex4
```

```
Hello from the thread, myid: 1, PID: 12347, TID:12349!  
Hello from the thread, myid: 0, PID: 12347, TID:12348!  
Hello from the thread, myid: 2, PID: 12347, TID:12350!  
Hello from the thread, myid: 3, PID: 12347, TID:12351!
```

Example 5, passing multiple arguments

```
1  struct pthread_args
2  {
3      long thread_id ;
4      long num_threads ;
5  };
6
7  void* kernel (void* args){
8      struct pthread_args *arg = (struct pthread_args*) args;
9      printf("Hello from the thread, number of threads: %ld, myid: %ld, PID: %d, TID:%d!\n", \
10 arg->num_threads, arg->thread_id, getpid(), (int) gettid());
11      return NULL;
12 }
```

Example 5, passing multiple arguments (cont.)

```
1  int main(int argc, char *argv[])
2  {
3      int num_threads=4;
4      pthread_t *threads = (pthread_t*) malloc (num_threads*sizeof(pthread_t));
5      struct pthread_args* args = (struct pthread_args*) malloc (num_threads*sizeof (struct pthread_args));
6
7      for (int i = 0; i < num_threads; ++i) {
8          //set the id and num threads in args for the threads
9          args[i].thread_id=i;
10         args[i].num_threads=num_threads;
11         //pass the args as argument to the threads
12         pthread_create(&threads[i], NULL, kernel, args+i); // passing args[i] to threads[i]
13     }
14
15     for (int i = 0; i < num_threads; ++i) {
16         pthread_join(threads[i], NULL);
17     }
18
19     free(threads); free(args);
20     return 0;
21 }
```


Example 6, how to get data out of threads

```

1  struct pthread_args
2  {
3      int in ;
4      int out ;
5  };
6
7  void* kernel_double (void* args){
8      struct pthread_args *arg = (struct pthread_args*) args;
9      arg->out = 2*arg->in;
10     return NULL;
11 }

```

Example 6, how to get data out of threads (cont.)

```
1  int main(int argc, char *argv[])
2  {
3      int num_threads=4;
4      pthread_t *threads = (pthread_t*) malloc (num_threads*sizeof(pthread_t));
5      struct pthread_args* args = (struct pthread_args*) malloc (num_threads*sizeof (struct pthread_args));
6
7      for (int i = 0; i < num_threads; ++i) {
8          args[i].in=i; //set the input in args
9          pthread_create(&threads[i], NULL, kernel_double, args+i);
10     }
11
12     for (int i = 0; i < num_threads; ++i) {
13         pthread_join(threads[i], NULL);
14     }
15
16     for (int i = 0; i < num_threads; ++i) {
17         printf("Double of %d is %d!\n", args[i].in, args[i].out);
18     }
19
20     free (threads); free (args); return 0;
21 }
```

Example 7, return data from threads

```
1 void* kernel_double (void* args){  
2     int in = *(int*) args;  
3     int *out = (int*) malloc (1*sizeof (int));  
4     *out = 2*in;  
5     return (void*)out;  
6 }
```

Example 7, return data from threads (cont.)

```
1  int main(int argc, char *argv[])
2  {
3      int num_threads=4;
4      pthread_t *threads = (pthread_t*) malloc (num_threads*sizeof(pthread_t));
5      int* in = (int*) malloc (num_threads*sizeof(int));
6
7      for (int i = 0; i < num_threads; ++i) {
8          in[i]=i; //set the input for the threads
9          pthread_create(&threads[i], NULL, kernel_double, in+i);
10     }
11
12     for (int i = 0; i < num_threads; ++i) {
13         int *out;
14         pthread_join(threads[i], (void*)&out);
15         printf("Double of %d is %d!\n", in[i], *out);
16         free(out);
17     }
18
19     free (threads); free (in); return 0;
20 }
```

What have we covered so far?

- Creating new threads with `pthread_create`
- Waiting for threads to finish with `pthread_join`
- Passing arguments to a pthread function
- Returning results from pthread function

Assignment 1: “Mandelbrot set” in parallel

Assignment: Mandelbrot


Starting this week, you have one week time.

- Use Pthreads to parallelize mandelbrot_draw()
- Your solution should have a speedup greater than 3.0 using 4 threads

```

1 void mandelbrot_draw( ... some args ) {
2     ...
3     for (int i = 0; i < y_resolution; i++)
4     {
5         for (int j = 0; j < x_resolution; j++)
6         {
7             //embarrassingly parrallel calculation of pixels
8             ...
9         }
10    }
11 }

```



Assignment: Mandelbrot (cont.)

Build the program

- Makefile:
make

Usage of the program

- Sequential:
`./mandelbrot_set_seq -h`
- Parallel:
`./mandelbrot_set_par -t 4 -r 480x380 -i 1000 -v [-2.0,0.5]x[-1.25,1.25] -f mandelbrot.ppm`

Assignment: Mandelbrot - provided files

- Makefile
 - contains rules to build executables
 - available targets: parallel, sequential, all (default), clean
 - 'mode=debug make [target]' to build debug version, use 'make clean' before
- main.c
 - main function - argument handling + file handling + call `draw_mandelbrot()`
- mandelbrot_set.h
 - Header file for `mandelbrot_set_*.c`
- mandelbrot.c
 - Defines helper functions
- mandelbrot_set_seq.c
 - Sequential version of `draw_mandelbrot()`
- student/mandelbrot_set_par.c
 - Implement the parallel version in this file

Assignment: Mandelbrot - provided files (cont.)

- `unit_test.c`
 - The unit tests that execute both the serial and parallel version to compare results.

Assignment: Extract, Build, and Run

1. Extract all files to the current directory
`tar -xvf assignment1.tar.gz`
2. Build the program
`make [sequential] [parallel] [unit_test]`
 - sequential: build the sequential program
 - parallel: build the parallel program
 - unit_test: builds the unit tests
3. Run the sequential program (with default parameters)
`student/mandelbrot_set_seq`
4. Run the parallel program (with 4 threads)
`student/mandelbrot_set_par -t 4`

Are you a windows user?

- Install linux in VirtualBox
 - Don't forget to assign multiple cores to the virtual machine
- Use the Machines at Rechnerhalle
 - Use Putty
 - ssh server: `lxhalle.informatik.tu-muenchen.de`
 - You need to get access from info point in informatik if you already don't have an account
- Ask the tutors; they will be more than happy to help you.

Submission

1. Log into the website
2. Go to Assignments
3. Use the link for Assignment 1
4. Upload your `mandelbrot_set_par.c` file
5. Press Submit

Parallel Programming
Home
Lecture material
Tutorial slides
Assignments
Schedule
Contacts
Impressum
Welcome ga59how Logout

Browse... No file selected.
Submit

Assignment Name	Submitted By	Date Of Submission	Status
Pthread 1; parallel Mandelbrot set	Amir Raoofy	April 16, 2018, 1 a.m.	✓ Accepted

Build step successful!
Correctness checks successful!
Pthread checks successful!
Runtime checks successful! - speedup: 3.511
Parallel Runtime: 0.88248

Top 10 Runtimes
0.88248